

regression-pr-02

January 4, 2024

```
[65]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[66]: df=pd.read_csv('../08-Linear-Regression-Models/Advertising.csv')
```

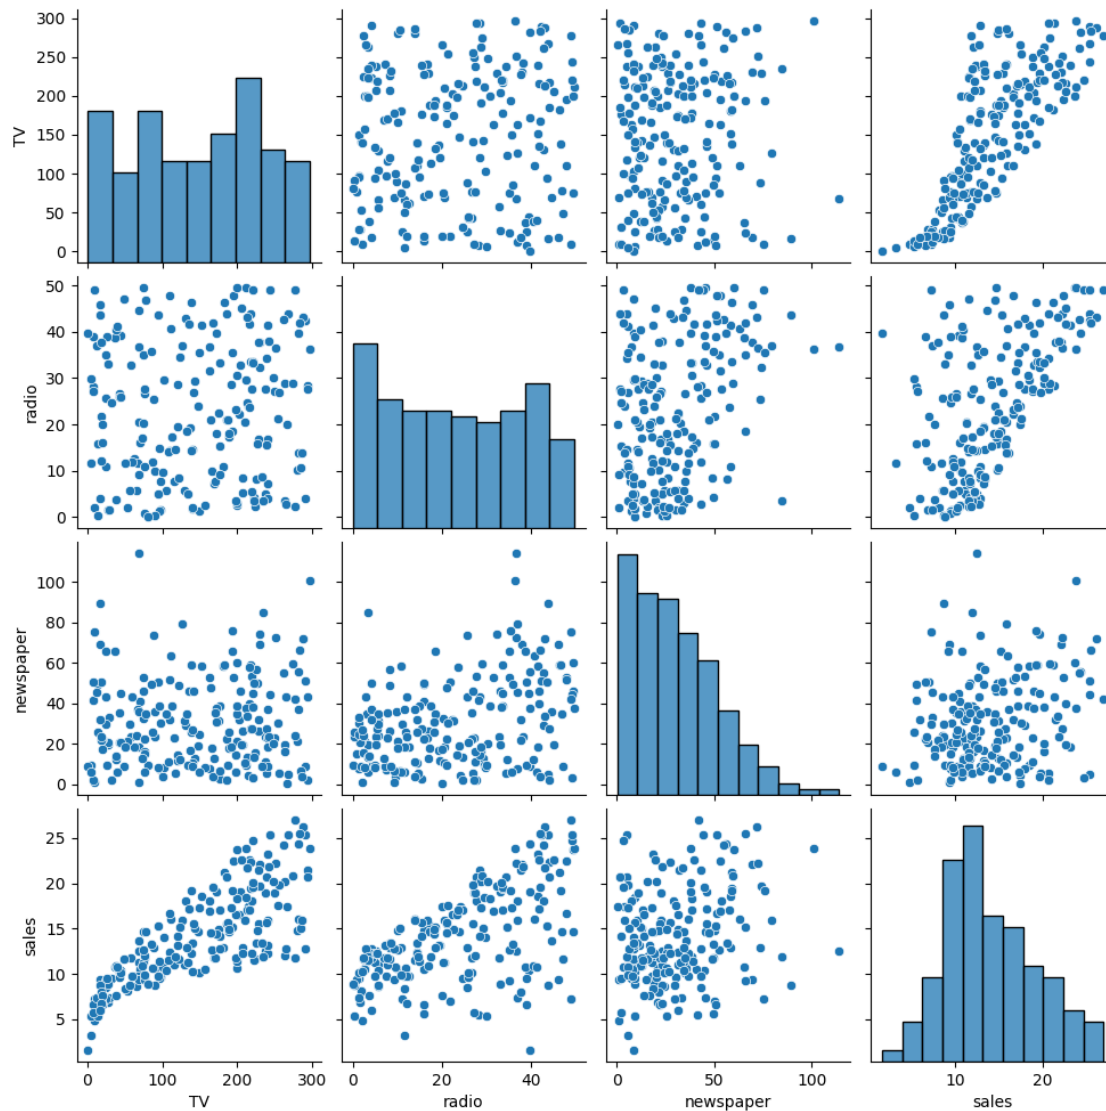
```
[67]: df.head()
```

```
[67]:
```

	TV	radio	newspaper	sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	9.3
3	151.5	41.3	58.5	18.5
4	180.8	10.8	58.4	12.9

```
[68]: sns.pairplot(df)
```

```
[68]: <seaborn.axisgrid.PairGrid at 0x7fc371d54520>
```



```
[69]: X=df.drop('sales',axis=1)
```

```
[70]: y=df['sales']
```

```
[71]: y
```

```
[71]: 0      22.1
      1      10.4
      2       9.3
      3      18.5
      4      12.9
      ...
     195     7.6
```

```
196      9.7
197     12.8
198     25.5
199     13.4
Name: sales, Length: 200, dtype: float64
```

```
[72]: from sklearn.model_selection import train_test_split
```

```
[73]: X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.
      ↪2,random_state=42)
```

```
[74]: len(df)
```

```
[74]: 200
```

```
[75]: len(X_train)
```

```
[75]: 160
```

```
[76]: from sklearn.linear_model import LinearRegression
```

```
[77]: model=LinearRegression()
```

```
[78]: model.fit(X_train,y_train)
```

```
[78]: LinearRegression()
```

```
[79]: test_pred=model.predict(X_test)
```

```
[80]: from sklearn.metrics import mean_absolute_error,mean_squared_error
```

```
[81]: df['sales'].mean()
```

```
[81]: 14.0225
```

```
[82]: mean_absolute_error(y_test,test_pred)
```

```
[82]: 1.4607567168117603
```


```
[83]: #RMSE
      np.sqrt(mean_squared_error(y_test,test_pred))
```

```
[83]: 1.78159966153345
```

```
[84]: test_residuals=y_test-test_pred
```


```
[85]: sns.scatterplot(x=y_test,y=test_residuals)
plt.axhline(y=0,color='r',ls='--')
```

```
[85]: <matplotlib.lines.Line2D at 0x7fc3902da5f0>
```




```
[86]: sns.kdeplot(test_residuals,fill=True)
```

```
[86]: <AxesSubplot:xlabel='sales', ylabel='Density'>
```



```
[87]: import scipy as sp
```

```
[88]: fig,ax=plt.subplots(figsize=(6,8))  
      _=sp.stats.probplot(test_residuals,plot=ax)
```



```
[89]: final_model=LinearRegression()
```

```
[90]: final_model.fit(X,y)
```

```
[90]: LinearRegression()
```

```
[91]: final_model.coef_
```

```
[91]: array([ 0.04576465,  0.18853002, -0.00103749])
```

```
[92]: y_hat=final_model.predict(X)
```

```
[93]: from joblib import dump,load
```

```
[94]: dump(final_model,'final_sales_model.joblib')
```

```
[94]: ['final_sales_model.joblib']
```

```
[ ]:
```

```
[95]: loaded_model=load('final_sales_model.joblib')
```

```
[96]: loaded_model
```

```
[96]: LinearRegression()
```

```
[97]: campaign=[[149,22,12]]
```

```
[98]: loaded_model.predict(campaign)
```

```
/home/mustafa/Desktop/ML/first/lib/python3.10/site-packages/sklearn/base.py:465:  
UserWarning: X does not have valid feature names, but LinearRegression was  
fitted with feature names  
warnings.warn(  

```

```
[98]: array([13.893032])
```

```
[ ]:
```