

# regression-pr-04

January 6, 2024

```
[1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[3]: df =pd.read_csv('../08-Linear-Regression-Models/Advertising.csv')
```

```
[6]: df.head()
```

```
[6]:
```

|   | TV    | radio | newspaper | sales |
|---|-------|-------|-----------|-------|
| 0 | 230.1 | 37.8  | 69.2      | 22.1  |
| 1 | 44.5  | 39.3  | 45.1      | 10.4  |
| 2 | 17.2  | 45.9  | 69.3      | 9.3   |
| 3 | 151.5 | 41.3  | 58.5      | 18.5  |
| 4 | 180.8 | 10.8  | 58.4      | 12.9  |

```
[7]: X=df.drop('sales',axis=1)
```

```
[8]: y=df['sales']
```

```
[9]: from sklearn.preprocessing import PolynomialFeatures
```

```
[10]: convertor=PolynomialFeatures(degree=3,include_bias=False)
```

```
[11]: poly_features=convertor.fit_transform(X)
```

```
[12]: from sklearn.model_selection import train_test_split
```

```
[14]: X_train,X_test,y_train,y_test=train_test_split(poly_features,y,test_size=0.
↪3,random_state=101)
```

```
[15]: from sklearn.preprocessing import StandardScaler
```

```
[16]: scaler=StandardScaler()
```

```
[17]: scaler.fit(X_train)
```

```

[17]: StandardScaler()

[18]: X_train=scaler.transform(X_train)

[19]: X_test=scaler.transform(X_test)

[20]: from sklearn.linear_model import Ridge

[21]: ridge_model=Ridge(alpha=10)

[22]: ridge_model.fit(X_train,y_train)

[22]: Ridge(alpha=10)

[23]: test_predictions=ridge_model.predict(X_test)

[24]: from sklearn.metrics import mean_absolute_error,mean_squared_error

[25]: MAE=mean_absolute_error(y_test,test_predictions)

[26]: MSE=mean_squared_error(y_test,test_predictions)
      RMSE=np.sqrt(MSE)

[27]: from sklearn.linear_model import RidgeCV

[43]: ridge_cv_model=RidgeCV(alphas=(0.1,1.0,10),scoring='neg_mean_absolute_error')

[44]: ridge_cv_model.fit(X_train,y_train)

[44]: RidgeCV(alphas=(0.1, 1.0, 10), scoring='neg_mean_absolute_error')

[45]: ridge_cv_model.alpha_

[45]: 0.1

[40]: from sklearn.metrics import get_scorer_names

[42]: get_scorer_names()

[42]: ['accuracy',
      'adjusted_mutual_info_score',
      'adjusted_rand_score',
      'average_precision',
      'balanced_accuracy',
      'completeness_score',
      'explained_variance',
      'f1',

```

```
'f1_macro',
'f1_micro',
'f1_samples',
'f1_weighted',
'fowlkes_mallows_score',
'homogeneity_score',
'jaccard',
'jaccard_macro',
'jaccard_micro',
'jaccard_samples',
'jaccard_weighted',
'matthews_corrcoef',
'max_error',
'mutual_info_score',
'neg_brier_score',
'neg_log_loss',
'neg_mean_absolute_error',
'neg_mean_absolute_percentage_error',
'neg_mean_gamma_deviance',
'neg_mean_poisson_deviance',
'neg_mean_squared_error',
'neg_mean_squared_log_error',
'neg_median_absolute_error',
'neg_negative_likelihood_ratio',
'neg_root_mean_squared_error',
'normalized_mutual_info_score',
'positive_likelihood_ratio',
'precision',
'precision_macro',
'precision_micro',
'precision_samples',
'precision_weighted',
'r2',
'rand_score',
'recall',
'recall_macro',
'recall_micro',
'recall_samples',
'recall_weighted',
'roc_auc',
'roc_auc_ovo',
'roc_auc_ovo_weighted',
'roc_auc_ovr',
'roc_auc_ovr_weighted',
'top_k_accuracy',
'v_measure_score']
```

```
[46]: test_predictions=ridge_cv_model.predict(X_test)
```

```
[47]: MAE=mean_absolute_error(y_test,test_predictions)
      RMSE=np.sqrt(mean_squared_error(y_test,test_predictions))
```

```
[48]: ridge_cv_model.coef_
```

```
[48]: array([ 5.40769392,  0.5885865 ,  0.40390395, -6.18263924,  4.59607939,
          -1.18789654, -1.15200458,  0.57837796, -0.1261586 ,  2.5569777 ,
          -1.38900471,  0.86059434,  0.72219553, -0.26129256,  0.17870787,
           0.44353612, -0.21362436, -0.04622473, -0.06441449])
```

```
[49]: ridge_cv_model.best_score_
```


```
[49]: -0.3749223340292964
```

```
[68]: plt_data=df.copy()
      plt_data['total']=plt_data['TV']+plt_data['radio']+plt_data['newspaper']
      test_data=np.linspace(0,120,60).reshape(20,3)
      test_data_sum=test_data.sum(axis=1)
      test_data=poly_features=convertor.fit_transform(test_data)

      test_data=scaler.transform(test_data)
      test_data
      predictions=ridge_cv_model.predict(test_data)
```

```
[69]: sns.scatterplot(data=plt_data,x='total',y='sales')
      plt.plot(test_data_sum,predictions,color='red')
```

```
[69]: [ <matplotlib.lines.Line2D at 0x7f00f28a7d90>]
```



[ ]: